

1.	Testing Environment	2
1.1	Software	2
1.2	Hardware.....	2
2.	Functional test (9 pass,1 fail)	2
2.1	create --PASS.....	2
2.2	assemble --PASS	2
2.3	check mdstat --PASS (include resync checking).....	3
2.4	format FS --PASS (such as ocfs2)	3
2.5	rd/wr operations --PASS.....	4
2.6	nodes update --PASS	4
2.7	home-cluster update --PASS.....	5
2.8	examine array --PASS.....	6
2.9	grow and switch bitmap --PASS (between internal and clustered).....	7
2.10	add and re-add failure device --FAIL.....	10
	Summarize.....	10
	a. test under pure rc1 environment	10
	b. verify by following packages based on rc1.....	11
3.	Stress testing	14
	Single node test -> PASS	15
4.	Performance compare.....	15
4.1	dd compared between local raid1 and md.	15
4.1.1	dd read compare	15
4.1.2	dd write compare	16
4.1.3	dd rd/wr compare	16
4.2	FIO test compare	16
4.2.1	10G + 24 threads -> Read	16
4.2.2	10G + 24 threads -> write.....	17
4.2.3	10G + 24 threads -> rand read	17
4.2.4	10G + 24 threads -> rand write	17
4.2.5	10G + 24 threads -> rw mixed rd.....	17
4.3	IOzone compared with local raid1, CLVM and cluster-md	18
5.	Performance details	18
5.1	Local RAID1 performance.....	18
5.1.1	dd include of read, write and rd/wr.	18
5.1.2	FIO testing	19
5.1.3	IOzone, ext4 format on local raid1	31
5.2	cluster-md performance	32
5.2.1	dd include of read, write and rd/wr.	32
5.2.2	FIO testing,	33
5.2.3	IOzone format ocfs2 FS on cluster-md	45
5.3	IOzone ocfs2 FS on CLVM	45

1. Testing Environment

1.1 Software

Kernel: sles12-sp1-rc1 (3.12.48-1-default)

HA: sles12-sp1-rc1

1.2 Hardware

Three physical machines

one for iscsi server(installed SLES12 GM),

IP: 147.2.207.200

the others for cluster(all installed SLES12 SP1 RC1).

IP: 147.2.207.235

IP: 147.2.207.233

2. Functional test (9 pass, 1 fail)

add and re-add doesn't work stable, refer to 10th as follow.

Note: function testing is based on VMs.

Node1:

zl-rc1-144, 147.2.208.144

Node2:

zl-rc1-166, 147.2.208.166

Node3:

zl-rc1-169, 147.2.208.169

2.1 create --PASS

create local raid1 and cluster-md all work normally.

2.2 assemble --PASS

all nodes of cluster can assemble md0 normally. Local raid1 works well.

2.3 check mdstat --PASS (include resync checking)

/proc/mdstat showed the correct infos. --PASS

zl-rc1-144:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sda[0] sdb[1]

419136 blocks super 1.2 [2/2] [UU]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

2.4 format FS --PASS (such as ocfs2)

format ocfs2 FS on md0 normally, mount and umount work well.

zl-rc1-144:~ # mkfs.ocfs2 /dev/md0

mkfs.ocfs2 1.8.2

Cluster stack: pcmk

Cluster name: zl-ha

Stack Flags: 0x0

NOTE: Feature extended slot map may be enabled

Label:

Features: sparse extended-slotmap backup-super unwritten inline-data
strict-journal-super xattr indexed-dirs refcount discontig-bg

Block size: 1024 (10 bits)

Cluster size: 4096 (12 bits)

Volume size: 429195264 (104784 clusters) (419136 blocks)

Cluster groups: 14 (tail covers 4944 clusters, rest cover 7680 clusters)

Extent allocator size: 2097152 (1 groups)

Journal size: 16777216

Node slots: 2

Creating bitmaps: done

Initializing superblock: done

Writing system files: done

Writing superblock: done

Writing backup superblock: 0 block(s)

Formatting Journals: done

Growing extent allocator: done

Formatting slot map: done

Formatting quota files: done

Writing lost+found: done

mkfs.ocfs2 successful

2.5 rd/wr operations --PASS

```
zl-rc1-144:~ # dd if=/dev/zero of=/dev/md0 bs=1M count=200
200+0 records in
200+0 records out
209715200 bytes (210 MB) copied, 13.0319 s, 16.1 MB/s
zl-rc1-144:~ # dd if=/dev/md0 of=md-io-test bs=1M count=200
200+0 records in
200+0 records out
209715200 bytes (210 MB) copied, 1.57526 s, 133 MB/s
zl-rc1-144:~ # du md-io-test -lh
200M    md-io-test
zl-rc1-144:~ # mount /dev/md0 /mnt/
zl-rc1-144:~ # cd /mnt/
zl-rc1-144:/mnt # ls
lost+found
zl-rc1-144:/mnt # touch 1 2 3
zl-rc1-144:/mnt # ls
1  2  3  lost+found
zl-rc1-144:/mnt # echo 123 > 1; echo 456 > 2; echo 789 > 3
zl-rc1-144:/mnt # cat 1 2 3
123
456
789
zl-rc1-144:/mnt # cd
zl-rc1-144:~ # umount /mnt
zl-rc1-144:~ # mount /dev/md0 /mnt/
zl-rc1-144:~ # rm /mnt/
1      2      3      lost+found/
zl-rc1-144:~ # rm /mnt/1 /mnt/2 /mnt/3
zl-rc1-144:~ # umount /mnt
zl-rc1-144:~ # mount /dev/md0 /mnt/
zl-rc1-144:~ # ls /mnt/
lost+found
zl-rc1-144:~ # umount /mnt
```

2.6 nodes update --PASS

```
#mdadm --assemble md0 --nodes=3 --update=nodes /dev/sda /dev/sdb
zl-rc1-144:~ # mdadm -X /dev/sda
    Filename : /dev/sda
    Magic : 6d746962
    Version : 5
```

UUID : 5c040480:d452377c:ca294357:669953bc
Chunksize : 64 MB
Daemon : 5s flush period
Write Mode : Normal
Sync Size : 419136 (409.38 MiB 429.20 MB)
Cluster nodes : 3
Cluster name : zl-ha
Node Slot : 0
Events : 0
Events Cleared : 0
State : OK
Bitmap : 7 bits (chunks), 0 dirty (0.0%)
Node Slot : 1
Events : 0
Events Cleared : 0
State : OK
Bitmap : 7 bits (chunks), 0 dirty (0.0%)
Node Slot : 2
Events : 0
Events Cleared : 0
State : OK
Bitmap : 7 bits (chunks), 0 dirty (0.0%)

2.7 home-cluster update --PASS

```
# mdadm --assemble md0 --update=home-cluster --home-cluster=zl-ha /dev/sda
/dev/sdb
zl-rc1-144:~ # mdadm -D /dev/md0
/dev/md0:
  Version : 1.2
  Creation Time : Mon Sep 28 17:32:18 2015
  Raid Level : raid1
  Array Size : 419136 (409.38 MiB 429.20 MB)
  Used Dev Size : 419136 (409.38 MiB 429.20 MB)
  Raid Devices : 2
  Total Devices : 2
  Persistence : Superblock is persistent

Intent Bitmap : Internal

  Update Time : Mon Sep 28 17:32:18 2015
  State : clean
  Active Devices : 2
  Working Devices : 2
```

Failed Devices : 0
Spare Devices : 0

Name : zl-rc1-144:md0 (local to host zl-rc1-144)
Cluster Name : zl-ha
UUID : 5c040480:d452377c:ca294357:669953bc
Events : 0

Number	Major	Minor	RaidDevice	State
0	8	0	active sync	/dev/sda
1	8	16	active sync	/dev/sdb

2.8 examine array --PASS

zl-rc1-144:~ # mdadm --examine /dev/sda
/dev/sda:
Magic : a92b4efc
Version : 1.2
Feature Map : 0x101
Array UUID : 5c040480:d452377c:ca294357:669953bc
Name : zl-rc1-144:md0 (local to host zl-rc1-144)
Cluster Name : zl-ha
Creation Time : Mon Sep 28 17:32:18 2015
Raid Level : raid1
Raid Devices : 2

Avail Dev Size : 838318 (409.40 MiB 429.22 MB)
Array Size : 419136 (409.38 MiB 429.20 MB)
Used Dev Size : 838272 (409.38 MiB 429.20 MB)
Data Offset : 544 sectors
Super Offset : 8 sectors
Unused Space : before=464 sectors, after=46 sectors
State : clean
Device UUID : 95ccf63d:dba3140b:c462583e:2c9e551f

Internal Bitmap : 8 sectors from superblock
Update Time : Mon Sep 28 17:32:18 2015
Checksum : 1c78885d - correct
Events : 0

Device Role : Active device 0

Array State : AA ('A' == active, '.' == missing, 'R' == replacing)

2.9 grow and switch bitmap --PASS (between internal and clustered)

zl-rc1-169:~ # mdadm -D /dev/md0

/dev/md0:

Version : 1.2

Creation Time : Mon Sep 28 14:44:42 2015

Raid Level : raid1

Array Size : 419136 (409.38 MiB 429.20 MB)

Used Dev Size : 419136 (409.38 MiB 429.20 MB)

Raid Devices : 3

Total Devices : 3

Persistence : Superblock is persistent

Intent Bitmap : Internal

Update Time : Mon Sep 28 14:48:21 2015

State : clean

Active Devices : 3

Working Devices : 3

Failed Devices : 0

Spare Devices : 0

Name : zl-rc1-169:md0 (local to host zl-rc1-169)

Cluster Name : cluster

UUID : 8aa7d880:86ccaba4:e5cdbbc0:d35e389a

Events : 64

Number	Major	Minor	RaidDevice	State
0	8	0	0	active sync /dev/sda
3	8	16	1	active sync /dev/sdb
2	8	32	2	active sync /dev/sdc

zl-rc1-169:~ # mdadm --grow /dev/md0 --bitmap=none

zl-rc1-169:~ # mdadm -D /dev/md0

/dev/md0:

Version : 1.2

Creation Time : Mon Sep 28 14:44:42 2015

Raid Level : raid1
Array Size : 419136 (409.38 MiB 429.20 MB)
Used Dev Size : 419136 (409.38 MiB 429.20 MB)
Raid Devices : 3
Total Devices : 3
Persistence : Superblock is persistent

Update Time : Mon Sep 28 14:56:57 2015
State : clean
Active Devices : 3
Working Devices : 3
Failed Devices : 0
Spare Devices : 0

Name : zl-rc1-169:md0 (local to host zl-rc1-169)
UUID : 8aa7d880:86ccaba4:e5cdbbc0:d35e389a
Events : 65

Number	Major	Minor	RaidDevice	State
0	8	0	0	active sync /dev/sda
3	8	16	1	active sync /dev/sdb
2	8	32	2	active sync /dev/sdc

zl-rc1-169:~ # mdadm --grow /dev/md0 --bitmap=internal
zl-rc1-169:~ # mdadm -D /dev/md0
/dev/md0:

Version : 1.2
Creation Time : Mon Sep 28 14:44:42 2015
Raid Level : raid1
Array Size : 419136 (409.38 MiB 429.20 MB)
Used Dev Size : 419136 (409.38 MiB 429.20 MB)
Raid Devices : 3
Total Devices : 3
Persistence : Superblock is persistent

Intent Bitmap : Internal

Update Time : Mon Sep 28 14:57:12 2015
State : clean
Active Devices : 3

Working Devices : 3

Failed Devices : 0

Spare Devices : 0

Name : zl-rc1-169:md0 (local to host zl-rc1-169)

UUID : 8aa7d880:86ccaba4:e5cdbbc0:d35e389a

Events : 66

Number	Major	Minor	RaidDevice	State
0	8	0	active sync	/dev/sda
3	8	16	active sync	/dev/sdb
2	8	32	active sync	/dev/sdc

zl-rc1-169:~ # mdadm --grow /dev/md0 --bitmap=none
zl-rc1-169:~ # mdadm -D /dev/md0
/dev/md0:

Version : 1.2

Creation Time : Mon Sep 28 14:44:42 2015

Raid Level : raid1

Array Size : 419136 (409.38 MiB 429.20 MB)

Used Dev Size : 419136 (409.38 MiB 429.20 MB)

Raid Devices : 3

Total Devices : 3

Persistence : Superblock is persistent

Update Time : Mon Sep 28 14:57:23 2015

State : clean

Active Devices : 3

Working Devices : 3

Failed Devices : 0

Spare Devices : 0

Name : zl-rc1-169:md0 (local to host zl-rc1-169)

UUID : 8aa7d880:86ccaba4:e5cdbbc0:d35e389a

Events : 67

Number	Major	Minor	RaidDevice	State
0	8	0	active sync	/dev/sda
3	8	16	active sync	/dev/sdb
2	8	32	active sync	/dev/sdc

zl-rc1-169:~ # mdadm --grow /dev/md0 --bitmap=clustered
zl-rc1-169:~ # mdadm -D /dev/md0
/dev/md0:

Version : 1.2
Creation Time : Mon Sep 28 14:44:42 2015
Raid Level : raid1
Array Size : 419136 (409.38 MiB 429.20 MB)
Used Dev Size : 419136 (409.38 MiB 429.20 MB)
Raid Devices : 3
Total Devices : 3
Persistence : Superblock is persistent

Intent Bitmap : Internal

Update Time : Mon Sep 28 14:57:37 2015
State : clean
Active Devices : 3
Working Devices : 3
Failed Devices : 0
Spare Devices : 0

Name : zl-rc1-169:md0 (local to host zl-rc1-169)
Cluster Name : cluster
UUID : 8aa7d880:86ccaba4:e5cdbbc0:d35e389a
Events : 68

Number	Major	Minor	RaidDevice	State
0	8	0	active sync	/dev/sda
3	8	16	active sync	/dev/sdb
2	8	32	active sync	/dev/sdc

2.10 add and re-add failure device --FAIL

Summarize

add and re-add can work well on local raid1,
also work well on the follow situation; cluster-md only start on one node of cluster, the other nodes of cluster must stop md0.

a. test under pure rc1 environment

Totally pure SLES12 SP1 RC1 kernel and RC1 HA environment.
"add" only can pass at first time, always got failure once repeatly; re-add cannot work.

b. verify by following packages based on rc1.

package (download them from Guoqing.):

cluster-md-kmp-default-0.0.1_k3.12.48_1-6.1.x86_64.rpm

mdadm-3.3.1-7.1.x86_64.rpm

1). assemble MD on all nodes of cluster, "add" and "re-add" operation got failure after "set fail" and "remove" device.

Create md:

```
# mdadm --create md0 --bitmap=clustered --raid-devices=3 --level=mirror /dev/sda  
/dev/sdb /dev/sdc
```

... ... sync the /etc/mdadm.conf to all nodes.

```
# cat /etc/mdadm.conf
```

```
DEVICE /dev/sda /dev/sdb /dev/sdc
```

```
ARRAY /dev/md/md0 metadata=1.2 UUID=00fd4a8c:af0db6c9:f5088e20:0900048c
```

Node1:

a) assemble and start cluster-md

```
zl-rc1-144:~ # mdadm -As
```

```
mdadm: /dev/md/md0 has been started with 3 drives.
```

```
zl-rc1-144:~ # cat /proc/mdstat
```

```
Personalities : [raid1]
```

```
md0 : active raid1 sda[0] sdc[2] sdb[1]
```

```
419136 blocks super 1.2 [3/3] [UUU]
```

```
bitmap: 0/1 pages [0KB], 65536KB chunk
```

```
unused devices: <none>
```

b) when "set fail" occurred

```
zl-rc1-144:~ # cat /proc/mdstat
```

```
Personalities : [raid1]
```

```
md0 : active raid1 sda[0] sdc[2] sdb[1](F)
```

```
419136 blocks super 1.2 [3/2] [U_U]
```

```
bitmap: 0/1 pages [0KB], 65536KB chunk
```

```
unused devices: <none>
```

c) when "remove" operation executed on other node.

```
zl-rc1-144:~ # cat /proc/mdstat
```

```
Personalities : [raid1]
md0 : active raid1 sda[0] sdc[2]
      419136 blocks super 1.2 [3/2] [U_U]
      bitmap: 0/1 pages [0KB], 65536KB chunk
```

unused devices: <none>

```
d)  it's the status when "add" or "re-add"
zl-rc1-144:~ # cat /proc/mdstat
Personalities : [raid1]
md0 : active raid1 sda[0] sdc[2]
      419136 blocks super 1.2 [3/2] [U_U]
      bitmap: 0/1 pages [0KB], 65536KB chunk
```

unused devices: <none>

```
# mdadm --manage /dev/md0 --add /dev/sdb
mdadm: add new device failed for /dev/sdb as 3: Invalid argument
```

```
# mdadm --manage /dev/md0 --re-add /dev/sdb
mdadm: --re-add for /dev/sdb to /dev/md0 is not possible
```

2). only assemble md array on one node of cluster, such as "zl-rc1-169" node started md and "zl-rc1-144, zl-rc1-166" nodes all stopped; then "add" and "re-add" all works well.

```
node1 :
# mdadm --create md0 --bitmap=clustered --raid-devices=3 --level=mirror /dev/sda
/dev/sdb /dev/sdc
```

```
# mdadm -Ds
ARRAY                /dev/md/md0                metadata=1.2                name=zl-rc1-169:md0
UUID=00fd4a8c:af0db6c9:f5088e20:0900048c
```

```
# cat /etc/mdadm.conf
DEVICE /dev/sda /dev/sdb /dev/sdc
ARRAY /dev/md/md0 metadata=1.2 UUID=00fd4a8c:af0db6c9:f5088e20:0900048c
```

```
# mdadm -S md0
```

zl-rc1-166:~ # cat /proc/mdstat

Personalities : [raid1]

unused devices: <none>

node2 :

zl-rc1-144:~ # cat /proc/mdstat

Personalities : [raid1]

unused devices: <none>

node3 :

zl-rc1-169:~ # mdadm -As

mdadm: /dev/md/md0 has been started with 3 drives.

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sda[0] sdc[2] sdb[1]

419136 blocks super 1.2 [3/3] [UUU]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

zl-rc1-169:~ # mdadm --manage /dev/md0 --fail /dev/sda

mdadm: set /dev/sda faulty in /dev/md0

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sda[0](F) sdc[2] sdb[1]

419136 blocks super 1.2 [3/2] [_UU]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

zl-rc1-169:~ # mdadm --manage /dev/md0 --remove /dev/sda

mdadm: hot removed /dev/sda from /dev/md0

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sdc[2] sdb[1]

419136 blocks super 1.2 [3/2] [_UU]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

zl-rc1-169:~ # mdadm --manage /dev/md0 --re-add /dev/sda

mdadm: re-added /dev/sda

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sda[0] sdc[2] sdb[1]

419136 blocks super 1.2 [3/3] [UUU]
bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

zl-rc1-169:~ # mdadm --manage /dev/md0 --fail /dev/sdb

mdadm: set /dev/sdb faulty in /dev/md0

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sda[0] sdc[2] sdb[1](F)

419136 blocks super 1.2 [3/2] [U_U]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

zl-rc1-169:~ # mdadm --manage /dev/md0 --remove /dev/sdb

mdadm: hot removed /dev/sdb from /dev/md0

You have new mail in /var/mail/root

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sda[0] sdc[2]

419136 blocks super 1.2 [3/2] [U_U]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

zl-rc1-169:~ # mdadm --manage /dev/md0 --add /dev/sdb

mdadm: re-added /dev/sdb

zl-rc1-169:~ # cat /proc/mdstat

Personalities : [raid1]

md0 : active raid1 sdb[1] sda[0] sdc[2]

419136 blocks super 1.2 [3/3] [UUU]

bitmap: 0/1 pages [0KB], 65536KB chunk

unused devices: <none>

3). local RAID1 all pass with "add" and "re-add".

3.Stress testing

testsuit named "ocfs2-test", format ocfs2 FS on md device.

Single node test -> PASS

the following cases are listed.

```
# vim ./ocfs2-test/bin/single_run-WIP.sh
...
run_aiostress ${LOGDIR} ${DEVICE} ${MOUNTPOINT}
run_mkfs ${LOGDIR} ${DEVICE} ${MOUNTPOINT} ${CLUSTER_STACK}
${CLUSTER_NAME}
run_tunefs ${LOGDIR} ${DEVICE} ${MOUNTPOINT} ${CLUSTER_STACK}
${CLUSTER_NAME}
run_backup_super ${LOGDIR} ${DEVICE} ${MOUNTPOINT}
${CLUSTER_STACK} ${CLUSTER_NAME}
run_inline_data ${LOGDIR} ${DEVICE} ${MOUNTPOINT}
${CLUSTER_STACK} ${CLUSTER_NAME}
run_xattr_test ${LOGDIR} ${DEVICE} ${MOUNTPOINT} ${CLUSTER_STACK}
${CLUSTER_NAME}
run_reflink_test ${LOGDIR} ${DEVICE} ${MOUNTPOINT}
${CLUSTER_STACK} ${CLUSTER_NAME}
...
```

4. Performance compare

Three disks exposed by iscsi-target(ISC SI-server IP: 147.2.207.200) are used as raid1.

4.1 dd compared between local raid1 and md.

4.1.1 dd read compare

local raid1:

```
# dd iflag=direct,nonblock if=/dev/md0 of=/dev/null bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 480.757 s, 89.3 MB/s
```

Cluster-md:

```
# dd iflag=direct,nonblock if=/dev/md0 of=/dev/null bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 415.498 s, 103 MB/s
```

4.1.2 dd write compare

local raid1:

```
dd oflag=direct,nonblock if=/dev/zero of=/dev/md0 bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 1433.75 s, 30.0 MB/s
```

Cluster-md:

```
# dd oflag=direct,nonblock if=/dev/zero of=/dev/md0 bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 1502.29 s, 28.6 MB/s
```

4.1.3 dd rd/wr compare

local raid1:

```
# dd iflag=direct,nonblock oflag=direct,nonblock if=/dev/md0 of=/home/test/iotest
bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 641.374 s, 67.0 MB/s
```

Cluster-md:

```
dd iflag=direct,nonblock oflag=direct,nonblock if=/dev/md0 of=/home/test/iotest
bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 653.715 s, 65.7 MB/s
```

4.2 FIO test compare

4.2.1 10G + 24 threads -> Read

Local raid1:

```
read : io=245760MB, bw=114062KB/s, iops=1782, runt=2206338msec
```

Cluster-md:

read : io=245760MB, bw=111973KB/s, iops=1749, runt=2247486msec

4.2.2 10G + 24 threads -> write

Local raid1:

write: io=245760MB, bw=54994KB/s, iops=859, runt=4576126msec

Cluster-md:

write: io=245760MB, bw=55056KB/s, iops=860, runt=4570953msec

4.2.3 10G + 24 threads -> rand read

Local raid1:

read : io=245760MB, bw=114797KB/s, iops=1793, runt=2192204msec

Cluster-md:

read : io=245760MB, bw=114799KB/s, iops=1793, runt=2192170msec

4.2.4 10G + 24 threads -> rand write

Local raid1:

write: io=245760MB, bw=32192KB/s, iops=502, runt=7817456msec

Cluster-md:

write: io=245760MB, bw=32431KB/s, iops=506, runt=7759705msec

4.2.5 10G + 24 threads -> rw mixed rd

Local raid1:

read : io=171929MB, bw=71145KB/s, iops=1111, runt=2474596msec

... ..

write: io=73832MB, bw=30552KB/s, iops=477, runt=2474596msec

Cluster-md:

read : io=171929MB, bw=66869KB/s, iops=1044, runt=2632840msec

... ..

write: io=73832MB, bw=28716KB/s, iops=448, runt=2632840msec

4.3 IOzone compared with local raid1, CLVM and cluster-md

Please refer to the attached logs.

5. Performance details

5.1 Local RAID1 performance

5.1.1 dd include of read, write and rd/wr.

Read

```
# time dd iflag=direct,nonblock if=/dev/md0 of=/dev/null bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 480.757 s, 89.3 MB/s
```

```
real    8m0.759s
user    0m0.072s
sys     0m4.468s
```

Write

```
# time dd oflag=direct,nonblock if=/dev/zero of=/dev/md0 bs=1M count=40960
40960+0 records in
40960+0 records out
```

42949672960 bytes (43 GB) copied, 1433.75 s, 30.0 MB/s

```
real 23m53.747s
user 0m0.072s
sys 0m13.784s
```

Read/write

```
# time dd iflag=direct,nonblock oflag=direct,nonblock if=/dev/md0
of=/home/test/iotest bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 641.374 s, 67.0 MB/s
```

```
real 10m41.395s
user 0m0.128s
sys 0m10.672s
```

5.1.2 FIO testing

Download source code from [github](#).

Operation contains **"read"**, **"write"**, **"randread"**, **"randwrite"** and **"rw-mix-read"**, threads=(num of processor) x 2, size= 10G, 20G

1) 10G + 24 threads

Read

```
test-read-10G: (g=0): rw=read, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64
```

...

```
fio-2.1.10
```

Starting 24 processes

```
test-read-10G: (groupid=0, jobs=24): err= 0: pid=16444: Mon Sep 28 02:37:30 2015
```

```
read : io=245760MB, bw=114062KB/s, iops=1782, runt=2206338msec
```

```
slat (usec): min=3, max=69367, avg=14.44, stdev=85.75
```

```
clat (msec): min=11, max=54063, avg=776.60, stdev=765.73
```

```
lat (msec): min=11, max=54063, avg=776.61, stdev=765.73
```

clat percentiles (msec):
| 1.00th=[322], 5.00th=[392], 10.00th=[429], 20.00th=[537],
| 30.00th=[537], 40.00th=[553], 50.00th=[685], 60.00th=[848],
| 70.00th=[1074], 80.00th=[1074], 90.00th=[1106], 95.00th=[1139],
| 99.00th=[1352], 99.50th=[1385], 99.90th=[4113], 99.95th=[16712],
| 99.99th=[16712]
bw (KB /s): min= 36, max=62976, per=5.18%, avg=5904.14, stdev=2455.57
lat (msec) : 20=0.01%, 50=0.01%, 100=0.27%, 250=0.41%, 500=16.78%
lat (msec) : 750=39.19%, 1000=11.09%, 2000=32.08%, >=2000=0.16%
cpu : usr=0.03%, sys=0.16%, ctx=297422, majf=0, minf=13063
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=3932160/w=0/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=245760MB, aggrb=114061KB/s, minb=114061KB/s,
maxb=114061KB/s, mint=2206338msec, maxt=2206338msec

Disk stats (read/write):

md0: ios=3932000/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=97362/0, aggrmerge=1212938/0, aggrticks=70122342/0,
aggrin_queue=70122761, aggrutil=100.00%

sdc: ios=159405/0, merge=2334393/0, ticks=139475888/0, in_queue=139477292,
util=100.00%

sdd: ios=132682/0, merge=1304423/0, ticks=70891140/0, in_queue=70890992,
util=100.00%

sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

Write

bj-ha-5:/opt/zhilong # cat log-write-10G.txt

test-write-10G: (g=0): rw=write, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-write-10G: (groupid=0, jobs=24): err= 0: pid=17383: Mon Sep 28 03:55:26 2015
write: io=245760MB, bw=54994KB/s, iops=859, runt=4576126msec
slat (usec): min=5, max=2629.1K, avg=432.61, stdev=24711.00
clat (msec): min=4, max=8552, avg=1784.59, stdev=384.44
lat (msec): min=5, max=8552, avg=1785.03, stdev=384.98
clat percentiles (msec):
| 1.00th=[660], 5.00th=[1303], 10.00th=[1565], 20.00th=[1713],
| 30.00th=[1713], 40.00th=[1713], 50.00th=[1713], 60.00th=[1713],
| 70.00th=[1729], 80.00th=[1762], 90.00th=[2180], 95.00th=[2573],
| 99.00th=[3261], 99.50th=[3490], 99.90th=[3916], 99.95th=[4146],
| 99.99th=[5014]
bw (KB /s): min= 24, max=26875, per=4.41%, avg=2427.82, stdev=936.93
lat (msec) : 10=0.01%, 20=0.01%, 50=0.01%, 100=0.01%, 250=0.03%
lat (msec) : 500=0.21%, 750=0.87%, 1000=0.95%,
2000=83.68%, >=2000=14.24%
cpu : usr=0.03%, sys=0.08%, ctx=246358, majf=0, minf=737
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=0/w=3932160/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

WRITE: io=245760MB, aggrb=54993KB/s, minb=54993KB/s, maxb=54993KB/s,
mint=4576126msec, maxt=4576126msec

Disk stats (read/write):

md0: ios=228/3932130, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=63/167504, aggrmerge=55/2451456, aggrticks=5160/231404701,
aggrin_queue=231428037, aggrutil=100.00%
sdc: ios=149/250018, merge=106/3678377, ticks=11136/356874316,
in_queue=356919632, util=100.00%
sdd: ios=40/252495, merge=59/3675993, ticks=4344/337339788,
in_queue=337364480, util=100.00%
sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

Rand read

bj-ha-5:/opt/zhilong # cat log-randread-10G.txt

test-randread-10G: (g=0): rw=randread, bs=64K-64K/64K-64K/64K-64K,
ioengine=libaio, iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randread-10G: (groupid=0, jobs=24): err= 0: pid=18927: Mon Sep 28 04:33:39
2015

read : io=245760MB, bw=114797KB/s, iops=1793, runt=2192204msec

slat (usec): min=3, max=1043.2K, avg=13321.16, stdev=75910.29

clat (msec): min=1, max=3089, avg=841.14, stdev=173.65

lat (msec): min=1, max=3089, avg=854.46, stdev=192.02

clat percentiles (msec):

| 1.00th=[404], 5.00th=[445], 10.00th=[461], 20.00th=[865],

| 30.00th=[881], 40.00th=[881], 50.00th=[889], 60.00th=[889],

| 70.00th=[898], 80.00th=[906], 90.00th=[922], 95.00th=[971],

| 99.00th=[1270], 99.50th=[1352], 99.90th=[1582], 99.95th=[1729],

| 99.99th=[1991]

bw (KB /s): min= 73, max=36096, per=4.19%, avg=4807.12, stdev=886.65

lat (msec) : 2=0.01%, 4=0.01%, 10=0.02%, 20=0.03%, 50=0.08%

lat (msec) : 100=0.09%, 250=0.24%, 500=11.41%, 750=2.65%, 1000=81.50%

lat (msec) : 2000=3.99%, >=2000=0.01%

cpu : usr=0.03%, sys=0.19%, ctx=127342, majf=0, minf=13035

IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%

submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%

complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%

issued : total=r=3932160/w=0/d=0, short=r=0/w=0/d=0

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=245760MB, aggrb=114796KB/s, minb=114796KB/s,
maxb=114796KB/s, mint=2192204msec, maxt=2192204msec

Disk stats (read/write):

md0: ios=3932069/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=1299893/0, aggrmerge=9691/0, aggrticks=117282861/0,
aggrin_queue=117283254, aggrutil=90.03%
sdc: ios=1948709/0, merge=13481/0, ticks=170119768/0, in_queue=170120268,
util=89.08%
sdd: ios=1950970/0, merge=15593/0, ticks=181728816/0, in_queue=181729496,
util=90.03%
sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

Randwrite

bj-ha-5:/opt/zhilong # cat log-randwrite-10G.txt
test-randwrite-10G: (g=0): rw=randwrite, bs=64K-64K/64K-64K/64K-64K,
ioengine=libaio, iodepth=64
...
fio-2.1.10
Starting 24 processes

test-randwrite-10G: (groupid=0, jobs=24): err= 0: pid=19822: Mon Sep 28 06:45:38
2015

write: io=245760MB, bw=32192KB/s, iops=502, runt=7817456msec
slat (usec): min=9, max=55867K, avg=47656.99, stdev=475909.49
clat (msec): min=15, max=59733, avg=2998.55, stdev=3565.29
lat (msec): min=15, max=59733, avg=3046.21, stdev=3592.74
clat percentiles (msec):
| 1.00th=[1287], 5.00th=[1713], 10.00th=[1713], 20.00th=[1729],
| 30.00th=[1942], 40.00th=[2278], 50.00th=[2507], 60.00th=[2737],
| 70.00th=[3130], 80.00th=[3556], 90.00th=[4015], 95.00th=[4490],
| 99.00th=[12125], 99.50th=[16712], 99.90th=[16712], 99.95th=[16712],
| 99.99th=[16712]
bw (KB /s): min= 1, max=16220, per=4.85%, avg=1561.42, stdev=532.97
lat (msec) : 20=0.01%, 50=0.01%, 100=0.01%, 250=0.01%, 500=0.03%
lat (msec) : 750=0.02%, 1000=0.03%, 2000=31.07%, >=2000=68.84%
cpu : usr=0.02%, sys=0.08%, ctx=251489, majf=0, minf=708
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=0/w=3932160/d=0, short=r=0/w=0/d=0

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

WRITE: io=245760MB, aggrb=32191KB/s, minb=32191KB/s, maxb=32191KB/s,
mint=7817456msec, maxt=7817456msec

Disk stats (read/write):

md0: ios=342/3932072, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=79/2610013, aggrmerge=55/9895, aggrticks=16420/424011650,
aggrin_queue=424034412, aggrutil=100.00%

sdc: ios=231/3930295, merge=155/1947, ticks=12036/118635684,
in_queue=118649840, util=74.51%

sdd: ios=6/3899744, merge=10/27740, ticks=37224/1153399268,
in_queue=1153453396, util=100.00%

sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

Rw-mix-read

bj-ha-5:/opt/zhilong # cat log-randrw\ -rwmixread\=70-10G.txt

test-randrw: (g=0): rw=randrw, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randrw: (groupid=0, jobs=24): err= 0: pid=22737: Mon Sep 28 07:28:33 2015

read : io=171929MB, bw=71145KB/s, iops=1111, runt=2474596msec

slat (usec): min=4, max=50493K, avg=9637.71, stdev=136664.52

clat (msec): min=2, max=107322, avg=901.72, stdev=1903.76

lat (msec): min=2, max=107322, avg=911.36, stdev=1912.56

clat percentiles (msec):

| 1.00th=[330], 5.00th=[545], 10.00th=[562], 20.00th=[586],

| 30.00th=[611], 40.00th=[717], 50.00th=[832], 60.00th=[873],

| 70.00th=[906], 80.00th=[1004], 90.00th=[1221], 95.00th=[1450],

| 99.00th=[1745], 99.50th=[2311], 99.90th=[15139], 99.95th=[16712],

| 99.99th=[16712]

bw (KB /s): min= 1, max=15168, per=4.63%, avg=3293.50, stdev=957.45

write: io=73832MB, bw=30552KB/s, iops=477, runt=2474596msec

slat (usec): min=7, max=50537K, avg=27443.20, stdev=302626.50

clat (msec): min=6, max=107322, avg=1038.06, stdev=2125.45
lat (msec): min=6, max=107322, avg=1065.50, stdev=2163.58
clat percentiles (msec):
| 1.00th=[519], 5.00th=[570], 10.00th=[594], 20.00th=[791],
| 30.00th=[840], 40.00th=[865], 50.00th=[889], 60.00th=[914],
| 70.00th=[996], 80.00th=[1188], 90.00th=[1434], 95.00th=[1565],
| 99.00th=[1844], 99.50th=[3425], 99.90th=[15533], 99.95th=[16712],
| 99.99th=[16712]
bw (KB /s): min= 1, max= 7746, per=4.64%, avg=1416.33, stdev=518.28
lat (msec) : 4=0.01%, 10=0.01%, 20=0.02%, 50=0.05%, 100=0.05%
lat (msec) : 250=0.11%, 500=1.27%, 750=33.17%, 1000=42.25%, 2000=22.46%
lat (msec) : >=2000=0.62%
cpu : usr=0.03%, sys=0.19%, ctx=166504, majf=0, minf=731
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=2750856/w=1181304/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=171929MB, aggrb=71144KB/s, minb=71144KB/s, maxb=71144KB/s,
mint=2474596msec, maxt=2474596msec

WRITE: io=73832MB, aggrb=30551KB/s, minb=30551KB/s, maxb=30551KB/s,
mint=2474596msec, maxt=2474596msec

Disk stats (read/write):

md0: ios=2750989/1181257, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=913305/786114, aggrmerge=3428/1312, aggrticks=73819260/63246560,
aggrin_queue=137071953, aggrutil=91.11%

sdc: ios=1384370/1178211, merge=7328/2854, ticks=127597484/110271136,
in_queue=237877756, util=91.11%

sdd: ios=1355546/1180133, merge=2958/1084, ticks=93860296/79468544,
in_queue=173338104, util=86.47%

sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

2) 20G + 24 threads

Read

bj-ha-5:/opt/zhilong # cat log-read-20G*

test-read-20G: (g=0): rw=read, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio, iodepth=64

...

fio-2.1.10

Starting 24 processes

test-read-20G: (groupid=0, jobs=24): err= 0: pid=23591: Mon Sep 28 08:45:21 2015

read : io=491520MB, bw=114203KB/s, iops=1784, runt=4407223msec

slat (usec): min=3, max=12909, avg=14.64, stdev=11.06

clat (msec): min=11, max=7482, avg=806.57, stdev=205.17

lat (msec): min=11, max=7482, avg=806.59, stdev=205.17

clat percentiles (msec):

| 1.00th=[322], 5.00th=[392], 10.00th=[392], 20.00th=[685],

| 30.00th=[758], 40.00th=[832], 50.00th=[865], 60.00th=[889],

| 70.00th=[938], 80.00th=[971], 90.00th=[996], 95.00th=[1012],

| 99.00th=[1057], 99.50th=[1205], 99.90th=[1713], 99.95th=[1975],

| 99.99th=[2311]

bw (KB /s): min= 78, max=82944, per=4.55%, avg=5196.46, stdev=1559.70

lat (msec) : 20=0.01%, 50=0.02%, 100=0.01%, 250=0.09%, 500=12.71%

lat (msec) : 750=16.69%, 1000=63.48%, 2000=6.98%, >=2000=0.02%

cpu : usr=0.03%, sys=0.16%, ctx=655094, majf=0, minf=13060

IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%

submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%

complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%

issued : total=r=7864320/w=0/d=0, short=r=0/w=0/d=0

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=491520MB, aggrb=114202KB/s, minb=114202KB/s,
maxb=114202KB/s, mint=4407223msec, maxt=4407223msec

Disk stats (read/write):

md0: ios=7864320/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=216612/0, aggrmerge=2403853/0, aggrticks=153367650/0,
aggrin_queue=153368130, aggrutil=100.00%

sdc: ios=382104/0, merge=5490368/0, ticks=331003688/0, in_queue=331004856,
util=100.00%

sdd: ios=267733/0, merge=1721192/0, ticks=129099264/0, in_queue=129099536,
util=100.00%

sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

start read 20G testing

Mon Sep 28 07:31:53 EDT 2015

finish read 20G testing

Mon Sep 28 08:45:21 EDT 2015

Write

bj-ha-5:/opt/zhilong # cat log-write-20G.txt

test-write-20G: (g=0): rw=write, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-write-20G: (groupid=0, jobs=24): err= 0: pid=25288: Mon Sep 28 11:25:09 2015

write: io=491520MB, bw=53052KB/s, iops=828, runt=9487299msec

slat (usec): min=5, max=2990.6K, avg=447.26, stdev=25982.45

clat (msec): min=4, max=11870, avg=1850.93, stdev=449.22

lat (msec): min=7, max=11870, avg=1851.37, stdev=449.73

clat percentiles (msec):

| 1.00th=[644], 5.00th=[1188], 10.00th=[1582], 20.00th=[1713],
| 30.00th=[1713], 40.00th=[1713], 50.00th=[1729], 60.00th=[1762],
| 70.00th=[1844], 80.00th=[2008], 90.00th=[2343], 95.00th=[2802],
| 99.00th=[3490], 99.50th=[3687], 99.90th=[4178], 99.95th=[4293],
| 99.99th=[4817]

bw (KB /s): min= 21, max=28228, per=4.49%, avg=2384.04, stdev=977.75

lat (msec) : 10=0.01%, 20=0.01%, 50=0.01%, 100=0.01%, 250=0.03%

lat (msec) : 500=0.19%, 750=1.02%, 1000=1.40%,
2000=77.14%, >=2000=20.22%

cpu : usr=0.03%, sys=0.07%, ctx=488006, majf=0, minf=738

IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%

submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,

64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=0/w=7864320/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

WRITE: io=491520MB, aggrb=53051KB/s, minb=53051KB/s, maxb=53051KB/s,
mint=9487299msec, maxt=9487299msec

Disk stats (read/write):

md0: ios=228/7864309, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=46/334603, aggrmerge=36/4904912, aggrticks=12341/475926781,
aggrin_queue=475964705, aggrutil=100.00%

sdc: ios=120/501525, merge=108/7357593, ticks=22928/688069940,
in_queue=688142084, util=100.00%

sdd: ios=18/502285, merge=2/7357144, ticks=14096/739710404,
in_queue=739752032, util=100.00%

sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

bj-ha-5:/opt/zhilong # cat time-write-20G.txt

start write 20G testing

Mon Sep 28 08:47:01 EDT 2015

finish write 20G testing

Mon Sep 28 11:25:09 EDT 2015

Rand read

bj-ha-5:/opt/zhilong # cat log-randread-20G.txt

test-randread-20G: (g=0): rw=randread, bs=64K-64K/64K-64K/64K-64K,
ioengine=libaio, iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randread-20G: (groupid=0, jobs=24): err= 0: pid=32299: Mon Sep 28 16:52:01
2015

read : io=491520MB, bw=25798KB/s, iops=403, runt=19510211msec

slat (usec): min=4, max=5687.2K, avg=59470.63, stdev=352918.06

clat (msec): min=2, max=13149, avg=3749.69, stdev=1287.78

lat (msec): min=2, max=14024, avg=3809.16, stdev=1355.33
clat percentiles (msec):
| 1.00th=[668], 5.00th=[1336], 10.00th=[1795], 20.00th=[2409],
| 30.00th=[3425], 40.00th=[3818], 50.00th=[4015], 60.00th=[4228],
| 70.00th=[4424], 80.00th=[4686], 90.00th=[5080], 95.00th=[5538],
| 99.00th=[6456], 99.50th=[6783], 99.90th=[7767], 99.95th=[8291],
| 99.99th=[9372]
bw (KB /s): min= 9, max=52254, per=4.85%, avg=1251.88, stdev=780.12
lat (msec) : 4=0.01%, 10=0.01%, 20=0.01%, 50=0.02%, 100=0.02%
lat (msec) : 250=0.06%, 500=0.15%, 750=1.10%, 1000=0.72%, 2000=11.23%
lat (msec) : >=2000=86.67%
cpu : usr=0.01%, sys=0.05%, ctx=245542, majf=0, minf=13038
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=7864320/w=0/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=491520MB, aggrb=25797KB/s, minb=25797KB/s, maxb=25797KB/s,
mint=19510211msec, maxt=19510211msec

Disk stats (read/write):

md0: ios=7864146/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=2607704/0, aggrmerge=12591/0, aggrticks=1034380844/0,
aggrin_queue=1034381173, aggrutil=91.65%

sdc: ios=3912683/0, merge=13330/0, ticks=1366947264/0, in_queue=1366945836,
util=89.07%

sdd: ios=3910429/0, merge=24444/0, ticks=1736195268/0, in_queue=1736197684,
util=91.65%

sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%

bj-ha-5:/opt/zhilong # cat time-randread-20G.txt

start randread 20G testing

Mon Sep 28 11:26:49 EDT 2015

finish randread 20G testing

Mon Sep 28 16:52:01 EDT 2015

Rand write

bj-ha-5:/opt/zhilong # cat log-randwrite-20G.txt

test-randwrite-20G: (g=0): rw=randwrite, bs=64K-64K/64K-64K/64K-64K,
ioengine=libaio, iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randwrite-20G: (groupid=0, jobs=24): err= 0: pid=6752: Tue Sep 29 00:26:48
2015

write: io=491520MB, bw=18515KB/s, iops=289, runt=27184473msec

slat (usec): min=9, max=93593K, avg=82815.37, stdev=448274.00

clat (msec): min=13, max=143664, avg=5214.35, stdev=2980.79

lat (msec): min=13, max=143664, avg=5297.17, stdev=3001.45

clat percentiles (msec):

| 1.00th=[1713], 5.00th=[2311], 10.00th=[3130], 20.00th=[4080],

| 30.00th=[4490], 40.00th=[4752], 50.00th=[5014], 60.00th=[5276],

| 70.00th=[5669], 80.00th=[6128], 90.00th=[6849], 95.00th=[7767],

| 99.00th=[10421], 99.50th=[14091], 99.90th=[16712], 99.95th=[16712],

| 99.99th=[16712]

bw (KB /s): min= 1, max=19401, per=4.66%, avg=863.10, stdev=352.01

lat (msec) : 20=0.01%, 50=0.01%, 100=0.01%, 250=0.01%, 500=0.01%

lat (msec) : 750=0.01%, 1000=0.01%, 2000=3.44%, >=2000=96.54%

cpu : usr=0.01%, sys=0.04%, ctx=514615, majf=0, minf=693

IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%

submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%

complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%

issued : total=r=0/w=7864320/d=0, short=r=0/w=0/d=0

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

WRITE: io=491520MB, aggrbw=18514KB/s, minb=18514KB/s, maxb=18514KB/s,
mint=27184473msec, maxt=27184473msec

Disk stats (read/write):

md0: ios=228/7864196, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,

aggrios=52/5227820, aggrmerge=36/13762, aggrticks=62312/70116317,
aggrin_queue=70239010, aggrutil=100.00%
sdc: ios=140/7821461, merge=108/38538, ticks=92076/4075585868,
in_queue=4075850040, util=100.00%
sdd: ios=16/7862000, merge=2/2750, ticks=94860/429730380,
in_queue=429834288, util=70.85%
sde: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%
bj-ha-5:/opt/zhilong # cat time-randwrite-20G.txt
start randwrite 20G testing ##
Mon Sep 28 16:53:41 EDT 2015
finish randwrite 20G testing ##
Tue Sep 29 00:26:48 EDT 2015

rw-mix-read

downgrade the priority, and kill it due to not enough time,

5.1.3 IOzone, ext4 format on local raid1

```

[1] ./iozone -i 0 -i 1 -i 2 -i 3 -i 4 -i 5 -i 6 -i 7 -i 8 -f /home/test/iozone.tmpfile -Rab /test-iozone.xls -g 8G -n 4G -C
Iozone: Performance Test of File I/O
Version $Revision: 3.430 $
Compiled for 64 bit mode.
Build: linux

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CVR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
vangel Bojaxhi, Ben England, Vikentsi Lapa.

Run began: Tue Sep 29 04:21:08 2015

Excel chart generation enabled
Auto Mode
Using maximum file size of 8388608 kilobytes.
Using minimum file size of 4194304 kilobytes.
Command line used: ./iozone -i 0 -i 1 -i 2 -i 3 -i 4 -i 5 -i 6 -i 7 -i 8 -f /home/test/iozone.tmpfile -Rab /test-iozone.xls -g 8G -n 4G -C
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      kB  reclen      write      rewrite      read      reread      random      random      bkwd      record      stride
      read      read      write      read      read      read      read      read      read      read
4194304      64      2631700      4063515      8130449      8055329      7850882      4602944      7931318      8541605      8095549      3888500      3774905      8104059      6912089
4194304      128      2599403      3983341      7962427      8390085      7847139      4813498      7909036      9504687      8888777      4121449      3910667      7963681      8353658
4194304      256      2605620      4102546      7351191      8123519      7495679      4713755      7353034      8286729      7448225      3799405      3738858      7343032      7366144
4194304      512      2659410      4127037      7386768      8256734      7421918      4754325      7410444      8013864      7566070      3929414      3749239      7445631      7517212
4194304      1024      2686711      4151675      7475959      7484645      7646931      4793649      7730169      8181423      7587571      3365715      3367369      7519719      8367049
4194304      2048      2690025      4192727      7549945      8420691      7637599      4761044      7534863      8244623      7608149      3892342      3043410      7464504      7567163
4194304      4096      2684222      4139823      7614861      7932115      7631268      4532353      7565565      7578331      7626808      2935627      2892248      7585170      7608110
4194304      8192      2495810      3757091      5840486      5987021      6329800      4209509      5914576      5280476      5904510      2869334      2847507      5909734      5993079
4194304      16384      2428854      3617391      4993236      5005637      4992803      4042099      4972265      4326531      4978460      2857569      2823354      4956774      5123777
8388608      64      361560      367596      8472872      8997589      8290048      190662      8387248      9298451      8768118      340854      364664      8492152      8464022
8388608      128      362992      369763      8373352      8578784      8438819      235285      8395073      9707632      8885737      363633      370772      8421569      8425114
8388608      256      364848      367809      7710017      7899739      7697975      261614      7721422      8099792      7671850      366827      368049      7713229      7766530
8388608      512      359943      366954      7796907      8237903      7924733      289090      7771864      7977031      7915086      366577      365622      7784651      7820073
8388608      1024      363003      368856      7880713      8097721      8026892      301096      7927872      7966131      8039415      366661      365003      7925700      7980703
8388608      2048      361908      372107      7988906      8157085      8029667      318899      7953864      8147442      8135094      365680      366405      7961429      7980607
8388608      4096      362037      366475      7988108      8128946      8018819      338158      7969688      8254024      8053671      361749      362611      7890943      8261926
8388608      8192      363402      367754      6039793      6311448      6099541      341651      6086807      5166633      6087228      364657      365334      6056065      6086605
8388608      16384      359135      366593      5064627      5171388      5088690      354730      5133557      4302569      5102665      361872      366315      5105416      5095137

iozone test complete.
Excel output is below:

"Writer report"
"4" "0" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"4194304" 0 0 0 0 2631700 2599403 2605620 2659410 2686711 2690025 2684222 2495810 2428854
"8388608" 0 0 0 0 361560 362992 354848 359943 363003 361908 362037 363402 359135

"Re-writer report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
```

Please refer to the detail log named “local-raid1-iozone.txt”

5.2 cluster-md performance

5.2.1 dd include of read, write and rd/wr.

READ

```
# time dd iflag=direct,nonblock if=/dev/md0 of=/dev/null bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 415.498 s, 103 MB/s
```

```
real    6m55.526s
user    0m0.060s
sys     0m3.776s
```

Write

```
# time dd oflag=direct,nonblock if=/dev/zero of=/dev/md0 bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 1502.29 s, 28.6 MB/s
```

```
real    25m2.294s
user    0m0.044s
sys     0m10.068s
```

Read/Write

```
# time dd iflag=direct,nonblock oflag=direct,nonblock if=/dev/md0
of=/home/test/iotest bs=1M count=40960
40960+0 records in
40960+0 records out
42949672960 bytes (43 GB) copied, 653.715 s, 65.7 MB/s
```

```
real    10m53.731s
user    0m0.144s
```


sys 0m10.680s

5.2.2 FIO testing,

operation contains "read", "write", "randread", "randwrite" and "rw-mix-read", threads=(num of processor) x 2, size= 10G, 20G, 30G

1) 10G + 24 threads

Read

```
bj-ha-5:/opt/md-test # cat log-read-10G.txt
test-read: (g=0): rw=read, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64
```

...

fio-2.1.10

Starting 24 processes

```
test-read: (groupid=0, jobs=24): err= 0: pid=10009: Sun Sep 27 08:05:03 2015
read : io=245760MB, bw=111973KB/s, iops=1749, runt=2247486msec
slat (usec): min=3, max=3318, avg=11.85, stdev= 8.72
clat (msec): min=12, max=54856, avg=789.86, stdev=434.81
lat (msec): min=12, max=54856, avg=789.87, stdev=434.81
clat percentiles (msec):
| 1.00th=[ 322], 5.00th=[ 383], 10.00th=[ 465], 20.00th=[ 482],
| 30.00th=[ 537], 40.00th=[ 660], 50.00th=[ 783], 60.00th=[ 922],
| 70.00th=[ 1020], 80.00th=[ 1045], 90.00th=[ 1188], 95.00th=[ 1221],
| 99.00th=[ 1270], 99.50th=[ 1319], 99.90th=[ 2008], 99.95th=[ 2311],
| 99.99th=[ 7242]
bw (KB /s): min= 38, max=115712, per=5.10%, avg=5715.55, stdev=2993.05
lat (msec) : 20=0.01%, 50=0.16%, 100=0.41%, 250=0.23%, 500=22.17%
lat (msec) : 750=25.38%, 1000=18.25%, 2000=33.30%, >=2000=0.10%
cpu      : usr=0.02%, sys=0.13%, ctx=309249, majf=0, minf=13061
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete  : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued   : total=r=3932160/w=0/d=0, short=r=0/w=0/d=0
```

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=245760MB, aggrb=111973KB/s, minb=111973KB/s,
maxb=111973KB/s, mint=2247486msec, maxt=2247486msec

Disk stats (read/write):

md0: ios=3931962/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=153637/0, aggrmerge=1811960/0, aggrticks=109535188/0,
aggrin_queue=109535582, aggrutil=100.00%

sdc: ios=156969/0, merge=2132120/0, ticks=133679772/0, in_queue=133679816,
util=100.00%

sdd: ios=150305/0, merge=1491800/0, ticks=85390604/0, in_queue=85391348,
util=99.99%

bj-ha-5:/opt/md-test # cat log-read-10G-time.txt

start read 10G testing

Sun Sep 27 07:27:35 EDT 2015

finish read testing

Sun Sep 27 08:05:03 EDT 2015

Write

bj-ha-5:/opt/md-test # cat log-write-10G*

start write 10G testing

Sun Sep 27 08:06:43 EDT 2015

finish write testing

Sun Sep 27 09:22:55 EDT 2015

test-write: (g=0): rw=write, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-write: (groupid=0, jobs=24): err= 0: pid=11026: Sun Sep 27 09:22:55 2015

write: io=245760MB, bw=55056KB/s, iops=860, runt=4570953msec

slat (usec): min=5, max=2565.2K, avg=452.02, stdev=25179.68

clat (msec): min=2, max=14247, avg=1782.70, stdev=397.56

lat (msec): min=2, max=14247, avg=1783.15, stdev=398.12

```

clat percentiles (msec):
| 1.00th=[ 635], 5.00th=[ 1221], 10.00th=[ 1565], 20.00th=[ 1713],
| 30.00th=[ 1713], 40.00th=[ 1713], 50.00th=[ 1713], 60.00th=[ 1713],
| 70.00th=[ 1729], 80.00th=[ 1745], 90.00th=[ 2212], 95.00th=[ 2606],
| 99.00th=[ 3228], 99.50th=[ 3425], 99.90th=[ 3884], 99.95th=[ 3982],
| 99.99th=[ 7504]
bw (KB /s): min= 24, max=38740, per=4.45%, avg=2448.73, stdev=1000.74
lat (msec) : 4=0.01%, 10=0.01%, 20=0.01%, 50=0.01%, 100=0.01%
lat (msec) : 250=0.03%, 500=0.23%, 750=0.93%, 1000=1.14%, 2000=83.76%
lat (msec) : >=2000=13.88%
cpu      : usr=0.03%, sys=0.06%, ctx=245931, majf=0, minf=738
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete  : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued   : total=r=0/w=3932160/d=0, short=r=0/w=0/d=0
latency   : target=0, window=0, percentile=100.00%, depth=64

```

Run status group 0 (all jobs):

```

WRITE: io=245760MB, aggrb=55055KB/s, minb=55055KB/s, maxb=55055KB/s,
mint=4570953msec, maxt=4570953msec

```

Disk stats (read/write):

```

md0: ios=342/3932055, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=130/252012, aggrmerge=110/3676263, aggrticks=14862/346463136,
aggrin_queue=346536780, aggrutil=100.00%
sdc: ios=222/251837, merge=153/3676295, ticks=16816/324810208,
in_queue=324862116, util=100.00%
sdd: ios=38/252188, merge=67/3676231, ticks=12908/368116064,
in_queue=368211444, util=100.00%

```

rand read

```

bj-ha-5:/opt/md-test # cat log-randread-10G*
## start randread 10G testing ##
      Sun Sep 27 09:24:35 EDT 2015
## finish randread testing ##

```

Sun Sep 27 10:01:08 EDT 2015

test-randread: (g=0): rw=randread, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio, iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randread: (groupid=0, jobs=24): err= 0: pid=12837: Sun Sep 27 10:01:07 2015

read : io=245760MB, bw=114799KB/s, iops=1793, runt=2192170msec

slat (usec): min=4, max=1000.6K, avg=13323.72, stdev=75908.81

clat (usec): min=855, max=2765.7K, avg=841172.22, stdev=171555.84

lat (usec): min=881, max=2929.7K, avg=854496.35, stdev=190125.71

clat percentiles (msec):

| 1.00th=[408], 5.00th=[445], 10.00th=[461], 20.00th=[865],
| 30.00th=[881], 40.00th=[881], 50.00th=[889], 60.00th=[889],
| 70.00th=[898], 80.00th=[906], 90.00th=[922], 95.00th=[971],
| 99.00th=[1237], 99.50th=[1336], 99.90th=[1582], 99.95th=[1680],
| 99.99th=[1909]

bw (KB /s): min= 74, max=66816, per=4.19%, avg=4809.40, stdev=915.76

lat (usec) : 1000=0.01%

lat (msec) : 4=0.01%, 10=0.01%, 20=0.03%, 50=0.07%, 100=0.08%

lat (msec) : 250=0.27%, 500=11.25%, 750=2.61%, 1000=81.69%, 2000=3.98%

lat (msec) : >=2000=0.01%

cpu : usr=0.03%, sys=0.18%, ctx=126611, majf=0, minf=13035

IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%

submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%

complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%

issued : total=r=3932160/w=0/d=0, short=r=0/w=0/d=0

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=245760MB, aggrb=114798KB/s, minb=114798KB/s,
maxb=114798KB/s, mint=2192170msec, maxt=2192170msec

Disk stats (read/write):

md0: ios=3932160/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=1949821/0, aggrmerge=14565/0, aggrticks=174918046/0,
aggrin_queue=174917166, aggrutil=89.92%

sdc: ios=1951725/0, merge=16378/0, ticks=181188636/0, in_queue=181187852, util=89.92%
sdd: ios=1947918/0, merge=12752/0, ticks=168647456/0, in_queue=168646480, util=88.81%

rand write

```
bj-ha-5:/opt/md-test # cat log-randwrite-10G*
## start randwrite 10G testing ##
      Sun Sep 27 10:02:48 EDT 2015
## finish randwrite testing ##
      Sun Sep 27 12:12:08 EDT 2015
test-randwrite:      (g=0):      rw=randwrite,      bs=64K-64K/64K-64K/64K-64K,
ioengine=libaio, iodepth=64
...
fio-2.1.10
Starting 24 processes
```

```
test-randwrite: (groupid=0, jobs=24): err= 0: pid=13736: Sun Sep 27 12:12:08 2015
write: io=245760MB, bw=32431KB/s, iops=506, runt=7759705msec
  slat (usec): min=7, max=65966K, avg=47273.16, stdev=350413.40
  clat (msec): min=6, max=76582, avg=2974.12, stdev=2578.05
    lat (msec): min=6, max=76582, avg=3021.40, stdev=2598.15
  clat percentiles (msec):
    | 1.00th=[ 1303],  5.00th=[ 1713], 10.00th=[ 1713], 20.00th=[ 1827],
    | 30.00th=[ 2180], 40.00th=[ 2409], 50.00th=[ 2638], 60.00th=[ 2966],
    | 70.00th=[ 3326], 80.00th=[ 3687], 90.00th=[ 4146], 95.00th=[ 4621],
    | 99.00th=[ 6259], 99.50th=[11731], 99.90th=[16712], 99.95th=[16712],
    | 99.99th=[16712]
  bw (KB /s): min=   0, max=26368, per=4.60%, avg=1491.43, stdev=513.74
  lat (msec) : 10=0.01%, 20=0.01%, 50=0.01%, 100=0.01%, 250=0.01%
    lat      (msec)      :      500=0.02%,      750=0.02%,      1000=0.04%,
2000=23.35%, >=2000=76.55%
  cpu      : usr=0.02%, sys=0.06%, ctx=252496, majf=0, minf=709
  IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
    submit   : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
    complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
  issued   : total=r=0/w=3932160/d=0, short=r=0/w=0/d=0
  latency  : target=0, window=0, percentile=100.00%, depth=64
```

Run status group 0 (all jobs):

WRITE: io=245760MB, aggrb=32431KB/s, minb=32431KB/s, maxb=32431KB/s,
mint=7759705msec, maxt=7759705msec

Disk stats (read/write):

md0: ios=342/3932068, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=96/3915771, aggrmerge=82/14168, aggrticks=38308/633055666,
aggrin_queue=633110144, aggrutil=100.00%

sdc: ios=170/3899883, merge=100/27732, ticks=12832/1149929820,
in_queue=1149976416, util=100.00%

sdd: ios=23/3931660, merge=65/604, ticks=63784/116181512,
in_queue=116243872, util=73.10%

rw-mix-read

bj-ha-5:/opt/md-test # cat log-randrw\ -rwmixread\=70-10G.txt

test-randrw: (g=0): rw=randrw, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randrw: (groupid=0, jobs=24): err= 0: pid=16386: Sun Sep 27 12:57:42 2015

read : io=171929MB, bw=66869KB/s, iops=1044, runt=2632840msec

slat (usec): min=4, max=42828K, avg=9934.49, stdev=171992.68

clat (msec): min=1, max=44439, avg=959.18, stdev=1762.23

lat (msec): min=1, max=44439, avg=969.11, stdev=1771.46

clat percentiles (msec):

| 1.00th=[314], 5.00th=[537], 10.00th=[562], 20.00th=[586],

| 30.00th=[611], 40.00th=[742], 50.00th=[840], 60.00th=[873],

| 70.00th=[914], 80.00th=[1020], 90.00th=[1287], 95.00th=[1467],

| 99.00th=[2089], 99.50th=[5473], 99.90th=[16712], 99.95th=[16712],

| 99.99th=[16712]

bw (KB /s): min= 4, max=19456, per=4.87%, avg=3253.42, stdev=1001.66

write: io=73832MB, bw=28716KB/s, iops=448, runt=2632840msec

slat (usec): min=7, max=42809K, avg=29881.78, stdev=322523.42

clat (msec): min=7, max=44439, avg=1102.48, stdev=1930.27

lat (msec): min=7, max=44764, avg=1132.36, stdev=1961.98
clat percentiles (msec):
| 1.00th=[510], 5.00th=[570], 10.00th=[594], 20.00th=[791],
| 30.00th=[840], 40.00th=[865], 50.00th=[889], 60.00th=[922],
| 70.00th=[1020], 80.00th=[1221], 90.00th=[1434], 95.00th=[1582],
| 99.00th=[2212], 99.50th=[8094], 99.90th=[16712], 99.95th=[16712],
| 99.99th=[16712]
bw (KB /s): min= 4, max= 7787, per=4.88%, avg=1399.94, stdev=532.30
lat (msec) : 2=0.01%, 4=0.01%, 10=0.01%, 20=0.02%, 50=0.07%
lat (msec) : 100=0.05%, 250=0.14%, 500=1.50%, 750=31.91%, 1000=42.05%
lat (msec) : 2000=23.00%, >=2000=1.25%
cpu : usr=0.03%, sys=0.15%, ctx=166994, majf=0, minf=731
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=2750856/w=1181304/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

READ: io=171929MB, aggrb=66868KB/s, minb=66868KB/s, maxb=66868KB/s,
mint=2632840msec, maxt=2632840msec

WRITE: io=73832MB, aggrb=28715KB/s, minb=28715KB/s, maxb=28715KB/s,
mint=2632840msec, maxt=2632840msec

Disk stats (read/write):

md0: ios=2751175/1181294, merge=0/0, ticks=0/0, in_queue=0,
util=0.00%, aggrios=1369876/1179133,
aggrmerge=5280/2047, aggrticks=119884958/102992140, aggrin_queue=222885698,
aggrutil=90.77%

sdc: ios=1393493/1180027, merge=3615/1236, ticks=100117684/83900596,
in_queue=184026136, util=86.82%

sdd: ios=1346260/1178239, merge=6945/2859, ticks=139652232/122083684,
in_queue=261745260, util=90.77%

2) 20G + 24 threads

Read

```
bj-ha-5:/opt/md-test # cat log-read-20G*
## start read 20G testing ##
      Sun Sep 27 13:01:02 EDT 2015
## finish read testing ##
      Sun Sep 27 14:14:13 EDT 2015
test-read: (g=0):  rw=read,  bs=64K-64K/64K-64K/64K-64K,  ioengine=libaio,
iodepth=64
...
fio-2.1.10
Starting 24 processes
```

```
test-read: (groupid=0, jobs=24): err= 0: pid=17526: Sun Sep 27 14:14:13 2015
 read : io=491520MB, bw=114623KB/s, iops=1790, runt=4391069msec
  slat (usec): min=3, max=1678, avg=11.61, stdev= 7.98
  clat (msec): min=6, max=5967, avg=773.39, stdev=265.39
    lat (msec): min=6, max=5967, avg=773.40, stdev=265.39
  clat percentiles (msec):
    | 1.00th=[ 359], 5.00th=[ 482], 10.00th=[ 545], 20.00th=[ 570],
    | 30.00th=[ 570], 40.00th=[ 570], 50.00th=[ 603], 60.00th=[ 848],
    | 70.00th=[ 1106], 80.00th=[ 1106], 90.00th=[ 1106], 95.00th=[ 1123],
    | 99.00th=[ 1156], 99.50th=[ 1188], 99.90th=[ 1680], 99.95th=[ 1844],
    | 99.99th=[ 2409]
  bw (KB /s): min= 50, max=115968, per=5.08%, avg=5818.71, stdev=2013.23
  lat (msec) : 10=0.01%, 20=0.01%, 50=0.05%, 100=0.01%, 250=0.04%
    lat      (msec)      :    500=7.81%,    750=50.82%,    1000=7.25%,
2000=33.99%, >=2000=0.03%
  cpu      : usr=0.02%, sys=0.13%, ctx=558859, majf=0, minf=13085
  IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
    submit    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
    complete   : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
  issued   : total=r=7864320/w=0/d=0, short=r=0/w=0/d=0
  latency  : target=0, window=0, percentile=100.00%, depth=64
```


Run status group 0 (all jobs):

READ: io=491520MB, aggrb=114622KB/s, minb=114622KB/s,
maxb=114622KB/s, mint=4391069msec, maxt=4391069msec

Disk stats (read/write):

md0: ios=7864184/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=275850/0, aggrmerge=3655450/0, aggrticks=202788660/0,
aggrin_queue=202791622, aggrutil=100.00%

sdc: ios=301856/0, merge=4497783/0, ticks=265281564/0, in_queue=265286736,
util=100.00%

sdd: ios=249844/0, merge=2813117/0, ticks=140295756/0, in_queue=140296508,
util=100.00%

Write

bj-ha-5:/opt/md-test # cat log-write-20G*

start write 20G testing

Sun Sep 27 14:15:53 EDT 2015

finish write testing

Sun Sep 27 16:53:16 EDT 2015

test-write: (g=0): rw=write, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-write: (groupid=0, jobs=24): err= 0: pid=19034: Sun Sep 27 16:53:16 2015

write: io=491520MB, bw=53307KB/s, iops=832, runt=9441863msec

slat (usec): min=5, max=2951.8K, avg=420.14, stdev=24701.07

clat (msec): min=5, max=13264, avg=1842.14, stdev=426.60

lat (msec): min=5, max=13264, avg=1842.56, stdev=427.08

clat percentiles (msec):

| 1.00th=[627], 5.00th=[1237], 10.00th=[1614], 20.00th=[1713],

| 30.00th=[1713], 40.00th=[1729], 50.00th=[1729], 60.00th=[1762],

| 70.00th=[1844], 80.00th=[1991], 90.00th=[2311], 95.00th=[2737],

| 99.00th=[3392], 99.50th=[3621], 99.90th=[4015], 99.95th=[4178],

| 99.99th=[4555]

bw (KB /s): min= 21, max=39071, per=4.49%, avg=2393.26, stdev=1001.09

lat (msec) : 10=0.01%, 20=0.01%, 50=0.01%, 100=0.01%, 250=0.03%

lat (msec) : 500=0.21%, 750=1.12%, 1000=1.18%,

2000=78.35%, >=2000=19.11%

cpu : usr=0.02%, sys=0.06%, ctx=488248, majf=0, minf=738
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued : total=r=0/w=7864320/d=0, short=r=0/w=0/d=0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

WRITE: io=491520MB, aggrb=53306KB/s, minb=53306KB/s, maxb=53306KB/s,
mint=9441863msec, maxt=9441863msec

Disk stats (read/write):

md0: ios=342/7864288, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=96/501910, aggrmerge=82/7357419, aggrticks=12914/715405288,
aggrin_queue=715469476, aggrutil=100.00%
sdc: ios=155/504445, merge=100/7354826, ticks=7080/670239748,
in_queue=670285076, util=100.00%
sdd: ios=38/499376, merge=65/7360013, ticks=18748/760570828,
in_queue=760653876, util=100.00%

randread

bj-ha-5:/opt/md-test # cat log-randread-20G*

start randread 20G testing

Sun Sep 27 16:54:56 EDT 2015

finish randread testing

Sun Sep 27 22:20:57 EDT 2015

test-randread: (g=0): rw=randread, bs=64K-64K/64K-64K/64K-64K, ioengine=libaio,
iodepth=64

...

fio-2.1.10

Starting 24 processes

test-randread: (groupid=0, jobs=24): err= 0: pid=22344: Sun Sep 27 22:20:56 2015

read : io=491520MB, bw=25734KB/s, iops=402, runt=19558360msec

```

slat (usec): min=4, max=6715.2K, avg=59622.64, stdev=351241.15
clat (msec): min=2, max=14430, avg=3758.68, stdev=1200.36
lat (msec): min=2, max=14430, avg=3818.30, stdev=1270.15
clat percentiles (msec):
| 1.00th=[ 709], 5.00th=[ 1467], 10.00th=[ 1876], 20.00th=[ 2573],
| 30.00th=[ 3490], 40.00th=[ 3752], 50.00th=[ 3982], 60.00th=[ 4228],
| 70.00th=[ 4490], 80.00th=[ 4752], 90.00th=[ 5080], 95.00th=[ 5342],
| 99.00th=[ 5800], 99.50th=[ 6194], 99.90th=[ 7439], 99.95th=[ 7832],
| 99.99th=[ 8979]
bw (KB /s): min= 9, max=38730, per=4.73%, avg=1216.22, stdev=703.65
lat (msec) : 4=0.01%, 10=0.01%, 20=0.01%, 50=0.03%, 100=0.02%
lat (msec) : 250=0.06%, 500=0.15%, 750=0.91%, 1000=0.73%, 2000=10.24%
lat (msec) : >=2000=87.86%
cpu      : usr=0.01%, sys=0.04%, ctx=246275, majf=0, minf=13039
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%,
32=0.1%, >=64=100.0%
submit    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%
complete  : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%
issued    : total=r=7864320/w=0/d=0, short=r=0/w=0/d=0
latency    : target=0, window=0, percentile=100.00%, depth=64

```

Run status group 0 (all jobs):

```

READ: io=491520MB, aggrb=25734KB/s, minb=25734KB/s, maxb=25734KB/s,
mint=19558360msec, maxt=19558360msec

```

Disk stats (read/write):

```

md0: ios=7864320/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,
aggrios=3911329/0, aggrmerge=19125/0, aggrticks=1543710402/0,
aggrin_queue=1543712000, aggrutil=90.27%
sdc: ios=3909716/0, merge=22347/0, ticks=1551574116/0, in_queue=1551574980,
util=90.27%
sdd: ios=3912943/0, merge=15903/0, ticks=1535846688/0, in_queue=1535849020,
util=90.20%

```

randwrite

```

bj-ha-5:/opt/md-test # cat log-randwrite-20G*

```

start randwrite 20G testing

Sun Sep 27 22:22:37 EDT 2015

test-randwrite: (g=0): rw=randwrite, bs=64K-64K/64K-64K/64K-64K,
ioengine=libaio, iodepth=64

...

fio-2.1.10

Starting 24 processes

fio: terminating on signal 2

test-randwrite: (groupid=0, jobs=24): err= 0: pid=29005: Sun Sep 27 22:38:20 2015

write: io=16983MB, bw=18432KB/s, iops=288, runt=943478msec

slat (usec): min=11, max=81651K, avg=83287.29, stdev=1089145.02

clat (msec): min=2, max=188687, avg=5139.77, stdev=13031.19

lat (msec): min=257, max=188701, avg=5223.06, stdev=13143.76

clat percentiles (msec):

| 1.00th=[1713], 5.00th=[2089], 10.00th=[2376], 20.00th=[2802],

| 30.00th=[3359], 40.00th=[3916], 50.00th=[4146], 60.00th=[4359],

| 70.00th=[4686], 80.00th=[5014], 90.00th=[5407], 95.00th=[5997],

| 99.00th=[15795], 99.50th=[16712], 99.90th=[16712], 99.95th=[16712],

| 99.99th=[16712]

bw (KB /s): min= 0, max= 3456, per=5.65%, avg=1040.82, stdev=458.72

lat (msec) : 4=0.01%, 10=0.01%, 20=0.01%, 50=0.01%, 100=0.01%

lat (msec) : 250=0.01%, 500=0.01%, 750=0.01%, 1000=0.01%, 2000=3.45%

lat (msec) : >=2000=96.53%

cpu : usr=0.01%, sys=0.04%, ctx=25716, majf=0, minf=712

IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%, 32=0.3%, >=64=99.4%

submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.0%, >=64=0.0%

complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%,
64=0.1%, >=64=0.0%

issued : total=r=0/w=271722/d=0, short=r=0/w=0/d=0

latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):

WRITE: io=16983MB, aggrbw=18432KB/s, minbw=18432KB/s, maxbw=18432KB/s,
mint=943478msec, maxt=943478msec

Disk stats (read/write):

md0: ios=233/271722, merge=0/0, ticks=0/0, in_queue=0, util=0.00%,

aggrios=116/272170, aggrmerge=102/5, aggrticks=42380/70887820,
aggrin_queue=70930136, aggrutil=100.00%
sdc: ios=165/272221, merge=137/0, ticks=1648/14250316, in_queue=14252036,
util=73.58%
sdd: ios=68/272120, merge=67/11, ticks=83112/127525324, in_queue=127608236,
util=100.00%

rw-mix-read

downgrade the priority, and kill it due to not enough time,

5.2.3 IOzone format ocfs2 FS on cluster-md

Please refer to the Excel spreadsheet named "test-md-iozone"

5.3 IOzone ocfs2 FS on CLVM

Please refer to the TXT and xls file name "test-clvm-iozone"